

Graph Sparsification and Community Detection

A chronological review of the literature, 1996–2026,
with an audit of what this project holds and what it does not

Compiled 2026-07-27

Scope, and how to read the verification marks

This covers one question and its immediate neighbourhood: *does removing edges from a graph before running community detection help, hurt, or leave it unchanged?* It is not a survey of graph sparsification in general, nor of community detection in general. Works from those fields appear only where the pipeline depends on them.

Every entry carries a verification mark, because the value of a review of this kind is entirely in whether its statements can be trusted:

[FULL]

The paper was read end to end from the PDF, and any characterization of what it claims can be relied on. Eleven papers are in this class; reading notes for most are in `PAPER_RESTRUCTURE/references/`.

[META]

Bibliographic fields were read off a publisher record, DOI resolver, arXiv, or the Semantic Scholar API. Safe to cite. Says nothing about whether the paper was read.

[2ND]

Described from a fetched page or search summary, not read directly. Listed as a lead. *Not safe to characterize.*

Where a claim about a paper’s content appears below without a [FULL] mark, it is a statement about what the paper says it does, not a verified account of what it does.

1 Prehistory: what sparsification guaranteed before anyone applied it to communities

The theory arrived first and it never promised what the pipeline later assumed.

Benczúr and Karger (1996, journal 2015) [1] [META] preserve every cut of a graph to within $1 \pm \epsilon$ by sampling edges with probability proportional to connectivity and reweighting the survivors.

Spielman and Teng (2004, journal 2014) [33, 34] [META] and **Spielman and Srivastava (2008)** [32] [META] preserve the Laplacian quadratic form by sampling according to effective resistance [18, 8], giving a spectral sparsifier with $O(n \log n / \epsilon^2)$ edges.

What this does and does not license. The guarantee is about the Laplacian spectrum. Community structure is related to that spectrum through the eigenvectors that encode cuts [35, 28, 20], but it is not determined by it. Every later claim that a spectral sparsifier “preserves communities” is an extrapolation beyond the theorem, and the gap between the two is the space this literature occupies.

Two structural filters predate the pipeline and are still used as baselines: the *disparity filter* of **Serrano et al. (2009)** [26] [META], which keeps edges improbable under a local null, and its noise-corrected successor from **Coscia and Neffke (2017)** [6] [META].

2 The founding claim: Satuluri, Parthasarathy and Ruan, SIGMOD 2011

[FULL] [24]. This is the origin of the sparsify-then-detect pipeline and everything downstream is a response to it. Read in full; notes in `NOTES_satuluri2011.md`.

The method, *L-Spar*, ranks each edge by the Jaccard similarity of its endpoints’ neighbourhoods, computed via minhash signatures, and has each vertex retain its top $\lceil d_v^e \rceil$ edges. The abstract claims speedups “often in the range 10–50” with “little or no deterioration in the quality of the resulting clusters,” and, decisively, “for at least two of the four clustering algorithms, our sparsification consistently enables higher clustering accuracies.”

Which two. Transcribed from their Table 2 (p. 729), ground-truth F-score by algorithm: Metis 4W/0L, Graclus 3W/0L, Metis+MQI 2W/2L, MLR-MCL 1W/4L. The accuracy claim is about the two **fixed- k cut partitioners**. The flow-based algorithm is the one L-Spar does not help. This is the single most misquoted fact in the downstream literature.

What they got right, and what the field then lost. Their §4.3 states, verbatim, that clusters obtained from a sparsified graph cannot be scored on that same graph, “since that would tell us nothing about how well the sparsified graph retained the cluster structure in the original graph,” and they report every conductance on the original graph. They fix the number of communities once per network and give it to every arm, so granularity is matched by construction. They charge sparsification time to the reported speedup. On all three counts the 1996–2011 origin paper is stricter than much of what followed it.

Their own scope conditions. A synthetic sweep in their §4.5 reports the method beginning to outperform clustering on the unsparsified graph only from average degree ≈ 50 . Their speedups are measured against clustering runs lasting hours, with an approximate minhash similarity reported as two orders of magnitude cheaper than exact Jaccard. Their disclosed limitation is that L-Spar cannot handle bipartite or triangle-free structure without a bibliographic-coupling transform.

The two detectors the accuracy claim is about: **Graclus** [7] [META], multilevel weighted kernel k -means for normalized cut, with no balance constraint; and **Metis** [14] [META], multilevel k -way partitioning with a hard balance constraint. The flow-based one is **MLR-MCL** [23] [META]. All three are distributed by their authors and build on a current toolchain.

3 The parallel question: does sampling preserve community structure at all?

A separate lineage asks whether *any* reduction preserves communities, independent of whether it accelerates anything.

Leskovec and Faloutsos (2006) [16] [META] set the graph-sampling problem. **Maiya and Berger-Wolf (2010)** [19] [META] sample explicitly to preserve community structure.

Blagus et al. (2015) [2] [META] is the dissenting result and deserves more attention than it gets: sampling can *manufacture* the appearance of community structure rather than preserve it. Any measured “improvement” after reduction has to clear this before it can be called a gain.

Peng, Kolda and Pinar (2014) [22] [META] reduce by k -core before detecting, which removes vertices as well as edges and therefore scores its partition on a smaller, denser vertex set.

4 Adoption and extension, 2016–2021

Hamann et al. (2016) [11] [FULL] is the systematic comparison, and it released the implementations now distributed in standard graph libraries, including Local Degree, Local Similarity and K-Neighbor. Two of its findings matter here and are frequently ignored. First, verbatim: normalized mutual information “reports higher similarity values for a larger number of found communities ... This makes it unsuitable for comparing partitions on sparsified networks,” with the adjusted Rand index recommended instead. Second, that filtering “does lead the algorithm into finding different community structures,” and that the preserved structure “is not necessarily the same as the one the Louvain algorithm finds.”

Bravo-Hermsdorff and Gunderson (NeurIPS 2019) [3] [FULL] unify sparsification and coarsening, which had been treated as separate operations. Their observation is that for a planar graph edge deletion corresponds to edge contraction in the dual, and that with respect to Laplacian dynamics the two are the reciprocal limits of edge weight 0 and ∞ . Their algorithm reduces a graph while preserving the Laplacian pseudoinverse $L^\dagger$, unbiased and variance-minimizing. *It does not evaluate community detection*. Its comparisons are on preservation of $L^\dagger$ and spectral distance against other reduction algorithms. It cites Satuluri et al. but tests something else, so it is context for the pipeline rather than evidence about it.

Sotiropoulos and Tsourakakis (KDD 2021) [31] [FULL] introduce triangle-aware spectral sparsifiers. The paper is worth reading for its framing alone: it treats sparsification as an established and effective route to community detection whose *success requires explanation*, opening “how can we explain its empirical success?” By 2021 the pipeline is a premise, not a proposition.

Satuluri et al. (2020) [25] [META] put sparsification before community detection in production at Twitter. **Wu and Chen (2020)** [36] [META] train a GAN to produce the subgraph.

Laeuchli (2020) [15] [FULL] is the one analysis that identifies a regime where the approach provably pays: dense planted-partition graphs with spectral solvers, where detection cost is superlinear in the number of edges and the planted structure survives heavy deletion. It is the honest boundary of the pipeline’s scope, and it is not the regime the pipeline is usually applied to.

5 The GNN turn, and what it cost this question

Liu et al. (2023) [17] [FULL] introduce *DSpar*, approximating the effective resistance of (u, v) by $1/d_u + 1/d_v$ and thereby obtaining a spectral sparsifier at the cost of reading the degree sequence. It is designed for GNN training, not community detection, and its adoption for the latter is downstream extrapolation.

Hashemi et al. (IJCAI 2024) [12] [FULL] survey graph reduction across sparsification, coarsening and condensation. The most useful fact about this survey is negative and was established by text search: it contains four occurrences of “graph classification,” one of “node classification,” one of “clustering,” and **zero** of “community.” The modern reduction field has reoriented entirely around GNN prediction and has dropped community detection as an evaluation target. The question was asked in 2011–2016, answered optimistically, and then abandoned rather than settled.

6 The benchmark era: Chen et al., PVLDB 2024

[FULL] [4], notes in `NOTES_chen_and_fastcd.md`. Twelve sparsifiers, sixteen graph properties, fourteen networks. Conclusions: no single sparsifier preserves all properties; the choice should follow the downstream task; pruning fragments graphs; and fidelity to the full graph’s partition degrades monotonically with the prune rate.

Two things about this paper matter disproportionately.

What it owns. Fragmentation-with-pruning at benchmark scale, and monotone fidelity degradation, are established facts of this literature and belong to Chen et al.

What it does not measure. It never computes a quality objective for community detection. Its clustering metric is fidelity to the full graph’s partition, not whether the sparsified partition is *better or worse*. It also never times the downstream task. So the field’s two central premises, that sparsification preserves quality and delivers speed, are both untested in the field’s most rigorous benchmark. Further, its released code computes modularity for the sparsified condition *on the sparsified graph* (`metrics_nk.py:414`), which is the clearest available evidence that the scoring hazard is present in the tooling of the most careful work in the area.

7 Redundancy: why any of this is possible

A separate and largely disconnected strand explains why graphs can lose most of their edges at all.

The *metric backbone* of a weighted graph is the union of its all-pairs shortest paths; an edge outside it is redundant in the exact sense that no shortest path uses it. **Correia, Barrat and Rocha (2023)** [5] [META] report the backbone to be a small fraction of the edges of real contact networks while community structure is retained. **Drevetton, Chucuri, Grossglauser and Thiran (NeurIPS 2024)** [9] [FULL] prove

that deleting every edge outside the backbone preserves community structure on a broad class of weighted random graphs with communities.

Their experimental section deserves separate comment, because it is the most carefully controlled evaluation in this entire literature. They preprocess each graph into a proximity graph using weighted Jaccard similarity, convert to a distance graph, and compare the metric backbone against a threshold graph and against Spielman–Srivastava spectral sparsification *at equal edge count*, $|E^{mb}| = |E^\theta| = |E^{ss}|$, which is a matched-retention control. They evaluate with the *adjusted* Rand index against ground-truth communities, so chance is subtracted. They run three algorithms, a Bayesian block model via graph-tool, Leiden, and spectral clustering, and they include the unsparified graph as a baseline in every figure.

Their conclusion is preservation, not improvement: “the metric backbone and the spectral sparsification retain equally well the community structure across all data sets and for all clustering algorithms tested.” Thresholding fares worse because it disconnects the graph.

One observation of theirs cuts against the similarity-based line. In their Figure 2 the edges present in the backbone but absent from the threshold graph “are mostly inter-community edges,” and the edges the threshold graph keeps but the backbone drops are mostly intra-community. The metric backbone *retains* the bridges that L-Spar and its descendants are designed to remove, and still preserves the community structure. No runtime or speedup is reported.

He, Drineas and Khanna (2025) [13] [FULL] prove that under strong clusterability, characterized by a structure ratio $\Upsilon(k) = \lambda_{k+1}/\rho_G(k)$, sampling $O(\gamma^2 n \log n / \epsilon^2)$ edges *uniformly at random* yields a sparsifier whose top $(n-k)$ eigenspace remains approximately orthogonal to the cluster indicators. Importance sampling is bypassed entirely.

They do run experiments, on stochastic block models with four clusters of 200 vertices, on a hierarchical block model with sixteen clusters, and on LFR benchmark graphs of 800 vertices with the mixing parameter varied. The measured quantity is the largest principal angle between the bottom- k eigenvectors of the sparsified graph and the true cluster indicators, which is a subspace-alignment error and not a community detection quality score, so the results do not speak directly to whether detection improves.

Two findings matter here. Uniform sampling tracks effective-resistance sampling closely, and verbatim: “even in the poorly-clustered setting, uniform sampling still followed an error trajectory similar to that of effective resistance sampling.” And on well-clustered graphs uniform sampling was *better*, which they attribute, as a hypothesis they do not test, to uniform sampling being “biased towards undersampling cross cluster edges.”

If this survives contact with an actual detection experiment it is a significant negative for the entire similarity-based line, because the expensive selection rule would be unnecessary exactly where structure is strong, and where structure is weak the guarantee lapses anyway.

8 The current wave, 2024–2026

The pipeline is not historical. Five works in three years, all proposing or applying it.

Socievole and Pizzuti publish twice: ASONAM 2024 [29] [META] and *Soft Computing* 30:2109–2133, 2026 [30] [FULL], notes in NOTES_socievole_pizzuti_2026.md. The 2026 method weights a graph by effective resistance, applies a kernel, sparsifies by weight thresholding, and runs a genetic algorithm maximizing weighted modularity on the result. It is worth reading for three reasons. Its evaluation measure is NMI against ground truth rather than a self-scored objective, so it does not commit the scoring error. It compares against unsparified Louvain and Infomap throughout. And its own tables show those unsparified baselines matching or beating every sparsified variant it tests: at mixing 0.3 on 1000-node LFR, Louvain reaches NMI 1 against the sparsified method’s 0.601, and its §6.2 states that on the densest network “OmeGANet achieves only 0.3522 as NMI value, while Louvain and Infomap are able to perfectly match the ground-truth.” The gains it reports are of kernels and the genetic algorithm over its own no-kernel baseline, not of sparsification over no sparsification.

Setiadi, Yaakub and Abu Bakar (2025) [27] [FULL], k -core community-preserving sparsification, notes in NOTES_kcore_ijain2026.md. Reproduced and falsified in this project. Its reported modularity of 0.63 on Zachary’s karate club exceeds 0.4198, the maximum attainable on that graph over all partitions, which identifies the scored graph as the sparsified one.

Pari, Bhandari and Raha (2026) [21] [FULL], effective-resistance sparsification and community detection, notes in NOTES_effres_2026.md. Reports modularity 0.87 on the American College Football network,

whose maximum over all partitions is 0.6046. Same diagnosis.

Gottesbüren et al. (ESA 2025) [10] [META] come from graph partitioning rather than community detection and observe the cost point directly: sparsification overhead can cancel the speedup it is meant to produce.

9 What the corpus contains, and what it does not

Holdings

Papers read end to end	11
PDFs archived locally	11
Lines of verified reading notes	2,664
Bibliography entries, all fields from fetched records	72
On-subject works spanning 2011–2026	14

The eight read in full are Satuluri et al. 2011 (plus the author’s thesis), Hamann et al. 2016, Laeuchli 2020, Sotiropoulos and Tsourakakis 2021, Liu et al. 2023, Chen et al. 2024, Hashemi et al. 2024, Setiadi et al. 2025, Pari et al. 2026, and Socievole and Pizzuti 2026. Four of these were not merely read but *reproduced*, and two of the reproductions falsified the paper’s central claim.

Genuine gaps, in priority order

The three items that headed this list in the first version of this document have been closed: He, Drineas and Khanna, Bravo-Hermsdorff and Gunderson, and the experimental section of Drevet et al. were all read in full on 2026-07-27 and their entries above are now [FULL]. Reading them changed three things, recorded here rather than buried:

1. **An attribution error was found and corrected.** The NeurIPS 2019 unifying framework is by **Bravo-Hermsdorff and Gunderson**, not by Sotiropoulos and Tsourakakis. The wrong attribution entered this project from a line in `P1_EVIDENCE.md` that was explicitly marked [UNVERIFIED] and was then promoted into a bibliography entry without being checked. It is corrected in `refs.bib` with a comment recording what happened. No other [UNVERIFIED] line has been promoted, but the episode is the reason the verification marks in this document exist.
2. **Drevet et al. turn out to run the best-controlled evaluation in this literature**, not a weak one: matched edge counts across three sparsifiers, a chance-corrected index, three detectors, and the unsparsified graph present as a baseline throughout. Their claim is preservation, not improvement, and it should be cited as such.
3. **He et al. is more threatening to the similarity-based line than its abstract suggests**, because it does have experiments, including on LFR, and uniform sampling matches or beats effective-resistance sampling in all of them.

What remains open:

1. **Socievole and Pizzuti (ASONAM 2024) is unread.** [META] only. The 2026 *Soft Computing* paper is its extension and has been read, which limits the exposure.
2. **He et al. measure subspace alignment, not detection quality.** Nobody has run their uniform-sampling result through an actual community detection evaluation. That is an experiment, not a reading gap, and it is the most interesting one this review surfaces.
3. **The 2012–2015 window is thin** on this specific question. Repeated searches return the same core corpus, which is consistent with the Hashemi survey’s finding that the question went dormant, but absence of search results is weak evidence.
4. **Coarsening and condensation applied to community detection** is covered only through the Hashemi survey, Peng et al., and now Bravo-Hermsdorff and Gunderson, none of which evaluates detection quality. If the pipeline has a modern descendant it is here.

Assessment

For the purpose of grounding a paper on this question the corpus is now close to complete. Every work that proposes or applies sparsify-then-detect and that a reviewer is likely to name has been located, and eleven have been read end to end, four of those reproduced, two of the reproductions falsifying the paper’s central claim.

Searching for further instances of the pipeline has reached saturation: the same dozen works return across many differently-phrased queries, in three separate sessions, which is itself consistent with the Hashemi survey’s finding that the modern reduction literature dropped community detection as a target. Further literature search is unlikely to return anything the corpus does not hold.

What the review leaves behind is not a reading task but an experimental one. Two questions in the sections above are now precisely posed and unanswered by anybody: whether uniform edge sampling, which He et al. prove suffices for subspace preservation and which beats effective-resistance sampling in their own figures, also suffices for community detection quality; and whether the metric backbone, which keeps inter-community bridges and still preserves structure, behaves the same way under a detection-quality evaluation rather than a preservation one. Neither can be settled by reading.

References

- [1] András A Benczúr and David R Karger. Randomized approximation schemes for cuts and flows in capacitated graphs. *SIAM Journal on Computing*, 44(2):290–319, 2015.
- [2] Neli Blagus, Lovro Šubelj, Gregor Weiss, and Marko Bajec. Sampling promotes community structure in social and information networks. *Physica A: Statistical Mechanics and its Applications*, 432:206–215, 2015.
- [3] Gecia Bravo-Hermsdorff and Lee M. Gunderson. A unifying framework for spectrum-preserving graph sparsification and coarsening. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- [4] Yuhan Chen, Haojie Ye, Sanketh Vedula, Alex Bronstein, Ronald Dreslinski, Trevor Mudge, and Nishil Talati. Demystifying graph sparsification algorithms in graph properties preservation. *Proceedings of the VLDB Endowment*, 17(3):427–440, 2024.
- [5] Rion Brattig Correia, Alain Barrat, and Luis M. Rocha. Contact networks have small metric backbones that maintain community structure and are primary transmission subgraphs. *PLOS Computational Biology*, 19(2):e1010854, 2023.
- [6] Michele Coscia and Frank M. H. Neffke. Network backboning with noisy data. In *2017 IEEE 33rd International Conference on Data Engineering (ICDE)*, pages 425–436, 2017.
- [7] Inderjit S. Dhillon, Yuqiang Guan, and Brian Kulis. Weighted graph cuts without eigenvectors: A multilevel approach. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 29(11):1944–1957, 2007.
- [8] Peter G Doyle and J Laurie Snell. *Random walks and electric networks*, volume 22. American Mathematical Society, 1984.
- [9] Maximilien Dreveton, Charbel Chucui, Matthias Grossglauser, and Patrick Thiran. Why the metric backbone preserves community structure. In *Advances in Neural Information Processing Systems 38 (NeurIPS)*, 2024. arXiv:2406.03852.
- [10] Lars Gottesbüren, Nikolai Maas, Dominik Rosch, Peter Sanders, and Daniel Seemaier. Linear-time multilevel graph partitioning via edge sparsification. In *33rd Annual European Symposium on Algorithms (ESA)*, 2025. arXiv:2504.17615.
- [11] Michael Hamann, Gerd Lindner, Henning Meyerhenke, Christian L. Staudt, and Dorothea Wagner. Structure-preserving sparsification methods for social networks. *Social Network Analysis and Mining*, 6(1):22, 2016.
- [12] Mohammad Hashemi, Shengbo Gong, Juntong Ni, Wenqi Fan, B. Aditya Prakash, and Wei Jin. A comprehensive survey on graph reduction: Sparsification, coarsening, and condensation. In *Proceedings of the 33rd International Joint Conference on Artificial Intelligence (IJCAI), Survey Track*, 2024.
- [13] Kaiwen He, Petros Drineas, and Rajiv Khanna. Structure-aware spectral sparsification via uniform edge sampling. *arXiv preprint arXiv:2510.12669*, 2025.
- [14] George Karypis and Vipin Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing*, 20(1):359–392, 1998.
- [15] Jesse Laeuchli. Fast community detection with graph sparsification. In *Advances in Knowledge Discovery and Data Mining (PAKDD)*, volume 12084 of *Lecture Notes in Computer Science*, pages 291–304, 2020.

- [16] Jure Leskovec and Christos Faloutsos. Sampling from large graphs. In *Proceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 631–636, 2006.
- [17] Zirui Liu, Kaixiong Zhou, Zhimeng Jiang, Li Li, Rui Chen, Soo-Hyun Choi, and Xia Hu. Dspar: An embarrassingly simple strategy for efficient GNN training and inference via degree-based sparsification. *Transactions on Machine Learning Research*, 2023.
- [18] László Lovász. Random walks on graphs: A survey. *Combinatorics, Paul Erdős Is Eighty*, 2:1–46, 1993.
- [19] Arun S. Maiya and Tanya Y. Berger-Wolf. Sampling community structure. In *Proceedings of the 19th International Conference on World Wide Web (WWW)*, pages 701–710, 2010.
- [20] Mark E J Newman. Spectral methods for network community detection and graph partitioning. *arXiv preprint arXiv:1307.7729*, 2013.
- [21] Jayanta Pari, Pratibha Bhandari, and Soumyendu Raha. Effective resistance-based graph sparsification and community detection. *arXiv preprint arXiv:2606.26766*, 2026.
- [22] Chengbin Peng, Tamara G. Kolda, and Ali Pinar. Accelerating community detection by using k -core subgraphs. *arXiv preprint arXiv:1403.2226*, 2014.
- [23] Venu Satuluri and Srinivasan Parthasarathy. Scalable graph clustering using stochastic flows: Applications to community discovery. In *Proceedings of the 15th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 737–746, 2009.
- [24] Venu Satuluri, Srinivasan Parthasarathy, and Yiye Ruan. Local graph sparsification for scalable clustering. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data*, pages 721–732, 2011.
- [25] Venu Satuluri, Yao Wu, Xun Zheng, Yilei Qian, Brian Wichers, Qieyun Dai, Gui Ming Tang, Jerry Jiang, and Jimmy Lin. SimClusters: Community-based representations for heterogeneous recommendations at Twitter. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD)*, pages 3183–3193, 2020.
- [26] M. Ángeles Serrano, Marián Boguñá, and Alessandro Vespignani. Extracting the multiscale backbone of complex weighted networks. *Proceedings of the National Academy of Sciences*, 106(16):6483–6488, 2009.
- [27] Tedy Setiadi, Mohd Ridzwan Yaakub, and Azuraliza Abu Bakar. Community preserving sparsification based on K-core for enhanced community detection in attributed networks. *International Journal of Advances in Intelligent Informatics*, 11(4):550–566, 2025.
- [28] Jianbo Shi and Jitendra Malik. Normalized cuts and image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(8):888–905, 2000.
- [29] Annalisa Socievoli and Clara Pizzuti. Community detection in complex networks exploiting spectral graph sparsification for efficient disaster response. In *Advances in Social Networks Analysis and Mining (ASONAM)*, 2024.
- [30] Annalisa Socievoli and Clara Pizzuti. Effective resistance and kernel-based graph sparsification for community detection in complex networks. *Soft Computing*, 30:2109–2133, 2026.
- [31] Konstantinos Sotiropoulos and Charalampos E. Tsourakakis. Triangle-aware spectral sparsifiers and community detection. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, pages 1501–1509, 2021.
- [32] Daniel A Spielman and Nikhil Srivastava. Graph sparsification by effective resistances. In *Proceedings of the Fortieth Annual ACM Symposium on Theory of Computing*, pages 563–568, 2008.
- [33] Daniel A Spielman and Shang-Hua Teng. Nearly-linear time algorithms for graph partitioning, graph sparsification, and solving linear systems. In *Proceedings of the Thirty-Sixth Annual ACM Symposium on Theory of Computing*, pages 81–90, 2004.
- [34] Daniel A Spielman and Shang-Hua Teng. Nearly linear time algorithms for preconditioning and solving symmetric, diagonally dominant linear systems. *SIAM Journal on Matrix Analysis and Applications*, 35(3):835–885, 2014.
- [35] Ulrike von Luxburg. A tutorial on spectral clustering. *Statistics and Computing*, 17(4):395–416, 2007.

- [36] Hang-Yang Wu and Yi-Ling Chen. Graph sparsification with generative adversarial network. In *2020 IEEE International Conference on Data Mining (ICDM)*, 2020.